

Unit 3 Project - High-Performance Simulation Study (due 11-21)

Goal. Optimize your Unit 2 simulation study for computational efficiency and numerical stability. Apply some of the methods we have studied -- profiling, complexity theory, algorithmic improvements, array programming, and parallelization -- to reduce runtime.

Learning Objectives. By completing this project, you will:

- Profile code to identify computational bottlenecks
- Apply appropriate strategies to optimize your code

Project Requirements

Core Requirements

1. Document Baseline Performance

Before making any optimizations, document your Unit 2 simulation's baseline performance:

- **Runtime Profiling:** Use a profiler to identify where time is spent
- **Memory Profiling:** This step is optional
- **Computational Complexity Analysis:** Estimate how long a single simulation run takes as a function of your key parameter(s)
 - Example: If studying risk vs sample size n , time your code for various values of n
 - Assess computational complexity either by
 - Empirical analysis: e.g. fit a line on a log-log plot of runtime vs parameter value
 - Theoretically: an analysis of the main steps of your DGP(s), method(s), metric(s) as a function of the main parameter
- **Evidence of numerical instability:** Log any numerical warnings, convergence issues or exceptions raised during the execution of your simulation study that could be evidence of numerical instability
- **Baseline Metrics Document:** Create `docs/BASELINE.md` discussing:
 - Total runtime of entire simulation study (the version you submitted for Unit 2)
 - Results of a profiler (e.g. a table or a plot) summary showing main bottlenecks
 - Computational complexity of bottlenecks with theoretical and/or empirical evidence
 - Numerical warnings or convergence issues observed (document frequency and conditions)

2. Optimization Implementation

Implement strategies from **at least two** of the following categories:

Algorithmic Improvements:

- Replace inefficient algorithms with faster alternatives
- Eliminate redundant computations
- Use analytical solutions where possible instead of iterative methods

Array Programming:

- Vectorize operations to eliminate Python loops (even, potentially, loops over simulation replicates)
- Effectively leverage broadcasting

Numerical Stability:

- Identify and mitigate numerically unstable computations
- Implement stable algorithms for critical calculations (e.g., LogSumExp, stable variance)
- Add safeguards against overflow/underflow

Parallelization:

- Implement parallel simulation replicates
- Parallelize across simulation scenarios

Other approaches not enumerated here...

3. Optimization Documentation

Create `docs/OPTIMIZATION.md` that includes:

For each optimization implemented:

- **Problem identified:** What bottleneck or issue did profiling reveal?
- **Solution implemented:** What optimization strategy did you apply?
- **Code comparison:** Show before/after code snippets (can be brief)
- **Performance impact:** Runtime improvement, memory savings, or stability gains
- **Trade-offs:** Any costs (code complexity, readability, maintenance, precision)?

Profiling Evidence:

- Include profiler output or visualizations showing improvements
- Show flame graphs, timing comparisons, or memory plots if helpful

Lessons Learned: Answer any of the following:

- Which optimizations provided the best return on investment?
- What surprised you about where time was actually spent?
- Which optimizations were not worth the effort?

4. Updated Makefile

Extend your Makefile with new targets:

```
make profile      # Run profiling on representative simulation
make complexity   # Run computational complexity analysis (timing vs n)
make benchmark     # Run timing comparison: baseline vs optimized
make parallel       # Run optimized version with parallelization
make stability-check # Check for warnings/convergence issues across conditions
```

5. Performance Comparison Visualization

Create at least one visualization comparing baseline vs optimized performance:

- **Computational complexity plot:** Runtime vs key parameter (e.g., n) on log-log scale
 - Show both baseline and optimized versions
 - Include fitted lines or empirical complexity estimates
 - Consider plotting different components separately (DGP time, method time, etc.)
- **Overall timing comparison:** Show runtime improvements across different simulation scales or conditions
- **Speedup analysis:** Demonstrate scaling behavior (e.g., speedup vs number of cores for parallelization)

6. Regression Tests

Demonstrate that your optimizations preserve correctness:

- **Result validation:** Show that optimized code produces equivalent results to baseline. For example:
 - Use statistical tests comparing distributions of results
 - Check that summary statistics match within reasonable tolerance
 - Verify edge cases produce identical outputs
 - Create a test or script (`test_regression.py`) that verifies this.
- If your new code produces different results than your old code, explain why in `docs/OPTIMIZATION.md`.

7. Optional Extensions

Open-ended challenges for students who want to push further. Document your analysis in `docs/EXTENSIONS.md` if you choose to attempt one of these. You are also encouraged to explore some aspect of the efficiency/stability of your simulation study that I haven't mentioned,

and similarly document your findings in [docs/EXTENSIONS.md](#).

Variance Reduction Techniques

Explore statistical techniques, such as Rao-Blackwellization or antithetic variates, to reduce the number of simulation replications needed to achieve a desired level of accuracy.

Simulation Budget Optimization

Consider the trade-off between computational cost and statistical precision. For different values of some parameter, e.g. n , the metrics of interest will have different variances. Statistically, this implies you may be able to use a lower value of n_{sim} for larger values of n .

Computationally, this is advantageous, since large n scenarios are the slowest. Try to understand these relationships empirically and/or theoretically.

Pilot Study Design

Consider running a small-scale simulation study in order to get a sense of (1) how large n_{sim} should be or (2) what scenarios to prioritize. How would you design your pilot study? And how would the results inform your larger-scale study? Keep your main study aims in mind when navigating this.